

# Explainable Financial Alerts for Retail Investors

## Integrating Statistical Anomaly Detection and News–Market Impact Correlation

Henrique J. S. Santos

Advisor: Prof. Luís Gomes    Co-advisor: Rafael Silva

MEIA — Instituto Superior de Engenharia do Porto

Defence

# The problem

- Most Americans own stock; many now trade through commission-free apps.
- The US equity market reached **US\$62.2 trillion** (end 2024).
- Retail investors face *continuous* price moves and news, **without** the tools of banks and funds.
- The AI spreading through finance is largely **opaque** — a poor fit for someone who must act and live with the result.


## Two distinct challenges

- 1 Tell a *genuinely* unusual move from ordinary volatility.
- 2 Read what a news item might *mean*, from how similar events behaved before.

# What CLARION does — and does not do

It does *not* predict prices. It **explains**.

For each abrupt move or relevant news item, it shows *what* was detected, *how*, the *sources*, and which *past events* resemble it — with what happened to prices afterwards.




# Research questions

- RQ1** Can **transparent** statistical methods detect abrupt moves consistently enough to be useful, while staying explainable?
- RQ2** Can a news–market engine retrieve **genuinely analogous** precedents and quantify their impact **without lookahead**, better than trivial baselines?
- RQ3** Can the explanations be **faithful** to what the system computes, and useful to a retail investor?

## Contribution — Artificial Intelligence *Engineering*

Not a new algorithm: the **disciplined integration, application and critical evaluation** of existing components into a working, explainable, reproducible system, with a documented news-to-impact methodology.

# Architecture



Single-responsibility components, common schema; two triggers converge on one explanation step.

# Market trigger — transparent by design

Rolling z-score, no lookahead:

$$z_t = \frac{r_t - \mu_t}{\sigma_t}, \quad \text{flag if } |z_t| > k$$

$\mu_t, \sigma_t$  over the window *before*  $t$ . Every quantity is exposed to the explanation.

## Worked example (real)

TSLA, 24 Oct 2024 (post-earnings):  
 $\mu = -0.92\%$ ,  $\sigma = 2.73\%$ ,  $r = +19.8\%$   
 $\Rightarrow z = +7.61$  (flagged at  $k = 3$ ).

A fixed % threshold would treat a calm and a volatile stock alike; the z-score does not.

# News trigger — retrieve analogous precedents, measure impact

- 1 Embed the headline (Sentence-BERT).
- 2 Cosine similarity vs the knowledge base.
- 3 Return top- $k$  precedents.
- 4 Measure their impact at +1/ + 3/ + 5 days, from the event-day close (**no lookahead**).

*Case-based reasoning*: evidence, not a forecast.

## Worked example (real)

Query: “Nvidia demand surges on AI chip orders”


→ NVDA guidance (sim 0.60) +5d +3.5%

→ MSFT cloud-AI (sim 0.38) +5d +10.9%

→ NVDA accelerator (sim 0.38) +5d +4.9%

**Cross-ticker** match (MSFT); mean +5d  $\approx$  +6.5%.

# Result 1 — anomaly detection is consistent across the market




- Firing-rate range across 15 tickers: **0.015** (z-score) vs **0.344** (fixed %).
- $> 20\times$  more consistent — a **label-free** argument.
- Vs extreme-move proxy:  $F_1 = 0.516$  vs 0.218.
- Window ablation:  $F_1$  0.385 / 0.516 / 0.678 (10/20/60 d).

## Result 2 — semantic retrieval beats every baseline

| Method         | P@5   | P@10  |
|----------------|-------|-------|
| SBERT (MiniLM) | 0.514 | 0.478 |
| SBERT (MPNet)  | 0.538 | 0.506 |
| Lexical        | 0.346 | 0.314 |
| Recency        | 0.126 | 0.106 |
| Random         | 0.240 | 0.240 |

Cross-ticker, 3,714 headlines, 5 seeds ( $\pm \sim 0.01$ ).



Lift largest where vocabulary is distinctive (energy, health).

## Result 3 — explanations are faithful by construction

- Each alert is rendered *directly* from the computed objects.
- So the text **cannot diverge** from the computation.
- Checked by an automated test (alert reproduces each precedent's date/ticker/impact).
- Every alert ends with a no-prediction disclaimer.

### Real alert (news trigger)

```
News alert -- NVDA
"Nvidia raises guidance...AI data-centre"
Avg 1-day move:  -1.97% (over 5; 3 shown)
- MSFT (0.68) +1d -3.46%
- META (0.68) +1d -2.65%
- GOOGL (0.64) +1d -0.46%
observed past outcome, not a prediction
```

## Limitations (stated plainly)

- Relevance is a **same-sector proxy**, not a human judgement.
- Corpus is a **recent** free-tier window, not the multi-year FNSPID history.
- Impact is a **raw** post-event return, not market-adjusted (CAR).
- Similarity is **thematic, not directional**: a cluster can mix up/down moves, so the mean is theme-level *evidence*, not a forecast.
- **No human usefulness study** yet (RQ3 usefulness open).
- By design: **no price prediction, no trading** — profit metrics out of scope.

All results are reproducible from versioned scripts (re-run: identical numbers).

# Conclusions

- **RQ1 (yes):** transparent z-score detects consistently across the market and explains every flag.
- **RQ2 (yes, retrieval):** semantic retrieval finds analogous precedents far above baselines; impact measured without lookahead.
- **RQ3 (faithful yes; useful open):** explanations cannot diverge from the computation; human usefulness is future work.

## Contribution

A working, **explainable, reproducible** financial-alert system for retail investors — engineering existing components into a transparent whole, with an honest evaluation.

# CLARION

Explainable financial alerts for retail investors

Questions welcome.

# Anticipated questions

- **Why not predict prices?** Market efficiency (Fama 1970): public news is impounded fast; we *measure* reactions, honestly.
- **Why z-score, not Isolation Forest/GARCH?** Transparency for a low-dimensional signal explained to a non-expert; defensible simplicity.
- **Why precision@k and a sector proxy?** Standard ranked-retrieval metric; objective, reproducible; human relevance is future work.
- **Cross-ticker constraint?** An *evaluation* choice to avoid trivial self-matches; the live alert shows the most similar precedents regardless.
- **A positive headline retrieved negative precedents (−1.97%) — misleading?** Similarity matches *theme*, not direction; the mean is theme-level **evidence**, shown with every precedent and a no-prediction disclaimer. Future work: de-duplicate precedents and flag directional disagreement.
- **Is it reproducible?** Yes — fixed seeds, pinned deps, versioned scripts; re-running reproduces every number.